

Data Wrangling

with pandas Cheat Sheet
<http://pandas.pydata.org>

[Pandas API Reference](#) [Pandas User Guide](#)

Creating DataFrames

	a	b	c
1	4	7	10
2	5	8	11
3	6	9	12

```
df = pd.DataFrame(
    {"a": [4, 5, 6],
     "b": [7, 8, 9],
     "c": [10, 11, 12]},
    index = [1, 2, 3])
```

Specify values for each column.

```
df = pd.DataFrame(
    [[4, 7, 10],
     [5, 8, 11],
     [6, 9, 12]],
    index=[1, 2, 3],
    columns=['a', 'b', 'c'])
```

Specify values for each row.

		a	b	c
N	v			
D	1	4	7	10
	2	5	8	11
e	2	6	9	12

```
df = pd.DataFrame(
    {"a": [4, 5, 6],
     "b": [7, 8, 9],
     "c": [10, 11, 12]},
    index = pd.MultiIndex.from_tuples(
        [('d', 1), ('d', 2),
         ('e', 2)], names=['n', 'v']))
```

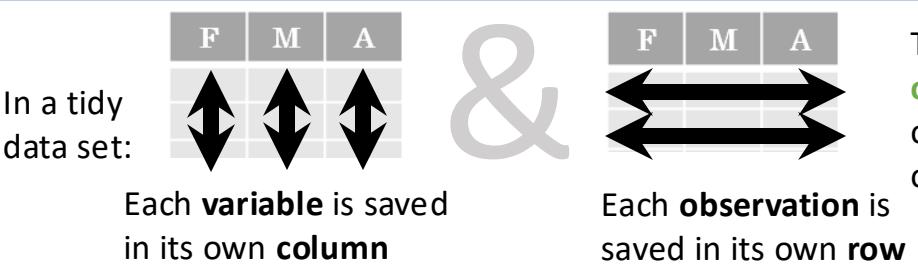
Create DataFrame with a MultiIndex

Method Chaining

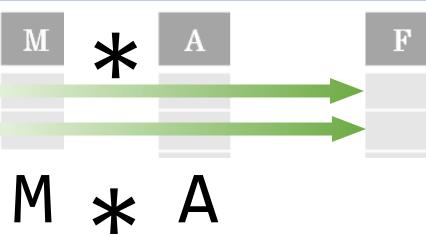
Most pandas methods return a DataFrame so that another pandas method can be applied to the result. This improves readability of code.

```
df = (pd.melt(df)
      .rename(columns={
          'variable':'var',
          'value':'val'})
      .query('val >= 200'))
```

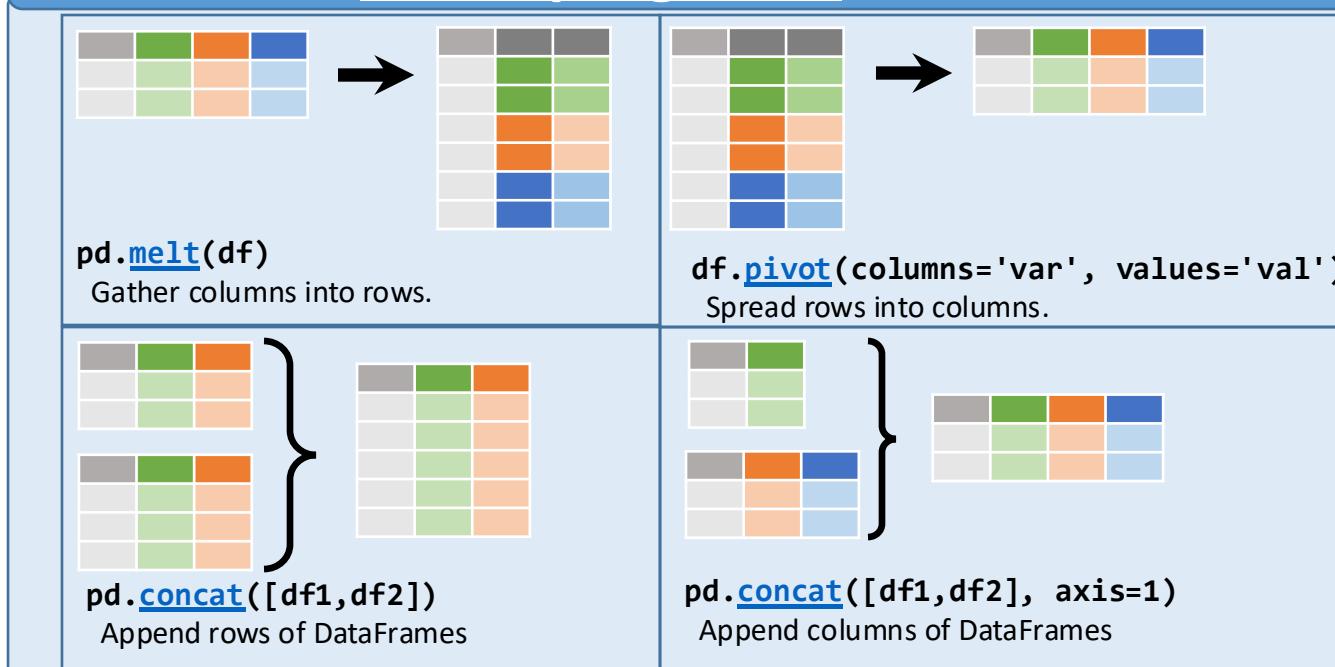
Tidy Data – A foundation for wrangling in pandas



Tidy data complements pandas's **vectorized operations**. pandas will automatically preserve observations as you manipulate variables. No other format works as intuitively with pandas.



Reshaping Data – Change layout, sorting, reindexing, renaming



df.sort_values('mpg')
Order rows by values of a column (low to high).

df.sort_values('mpg', ascending=False)
Order rows by values of a column (high to low).

df.rename(columns = {'y': 'year'})
Rename the columns of a DataFrame

df.sort_index()
Sort the index of a DataFrame

df.reset_index()
Reset index of DataFrame to row numbers, moving index to columns.

df.drop(columns=['Length', 'Height'])
Drop columns from DataFrame

Subset Observations - rows



df[df.Length > 7]
Extract rows that meet logical criteria.

df.drop_duplicates()
Remove duplicate rows (only considers columns).

df.sample(frac=0.5)
Randomly select fraction of rows.

df.sample(n=10)
Randomly select n rows.

df.nlargest(n, 'value')
Select and order top n entries.

df.nsmallest(n, 'value')
Select and order bottom n entries.

df.head(n)

Select first n rows.

df.tail(n)

Select last n rows.

Subset Variables - columns



df[['width', 'length', 'species']]
Select multiple columns with specific names.

df['width'] or df.width

Select single column with specific name.

df.filter(regex='regex')
Select columns whose name matches regular expression regex.

Using query

query() allows Boolean expressions for filtering rows.

df.query('Length > 7')

df.query('Length > 7 and Width < 8')

df.query('Name.str.startswith("abc")', engine="python")

Use **df.loc[]** and **df.iloc[]** to select only rows, only columns or both.

Use **df.at[]** and **df.iat[]** to access a single value by row and column.

First index selects rows, second index columns.

df.iloc[10:20]

Select rows 10-20.

df.iloc[:, [1, 2, 5]]

Select columns in positions 1, 2 and 5 (first column is 0).

df.loc[:, 'x2':'x4']

Select all columns between x2 and x4 (inclusive).

df.loc[df['a'] > 10, ['a', 'c']]

Select rows meeting logical condition, and only the specific columns.

df.iat[1, 2] Access single value by index

df.at[4, 'A'] Access single value by label

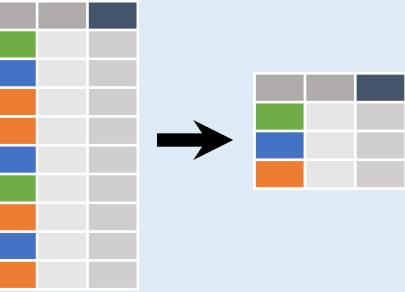
Logic in Python (and pandas)

<	Less than	!=	Not equal to
>	Greater than	df.column.isin(values)	Group membership
==	Equals	pd.isnull(obj)	Is NaN
<=	Less than or equals	pd.notnull(obj)	Is not NaN
>=	Greater than or equals	&, , ~, ^, df.any(), df.all()	Logical and, or, not, xor, any, all

regex (Regular Expressions) Examples

'.'	Matches strings containing a period '.'
'Length\$'	Matches strings ending with word 'Length'
'^Sepal'	Matches strings beginning with the word 'Sepal'
'^x[1-5]\$'	Matches strings beginning with 'x' and ending with 1,2,3,4,5
'^(?!Species\$).*''	Matches strings except the string 'Species'

Group Data



`df.groupby(by="col")`
Return a GroupBy object, grouped by values in column named "col".

`df.groupby(level="ind")`
Return a GroupBy object, grouped by values in index level named "ind".

The examples below can also be applied to groups. In this case, the function is applied on a per-group basis, and the returned vectors are of the length of the original DataFrame.

`shift(1)`
Copy with values shifted by 1.

`rank(method='dense')`

Ranks with no gaps.

`rank(method='min')`

Ranks. Ties get min rank.

`rank(pct=True)`

Ranks rescaled to interval [0, 1].

`rank(method='first')`

Ranks. Ties go to first value.

`shift(-1)`
Copy with values lagged by 1.

`cumsum()`

Cumulative sum.

`cummax()`

Cumulative max.

`cummin()`

Cumulative min.

`cumprod()`

Cumulative product.

All of the summary functions listed above can be applied to a group.

Additional GroupBy functions:

`size()`
Size of each group.

`agg(function)`
Aggregate group using function.

Summarize Data

`df['w'].value_counts()`
Count number of rows with each unique value of variable

`len(df)`
of rows in DataFrame.

`df.shape`
Tuple of # of rows, # of columns in DataFrame.

`df['w'].nunique()`
of distinct values in a column.

`df.describe()`
Basic descriptive and statistics for each column (or GroupBy).

`df.info()`
Prints a concise summary of the DataFrame.

`df.memory_usage()`
Prints the memory usage of each column in the DataFrame.

`df.dtypes()`
Prints a Series with the dtype of each column in the DataFrame.



pandas provides a large set of **summary functions** that operate on different kinds of pandas objects (DataFrame columns, Series, GroupBy, Expanding and Rolling (see below)) and produce single values for each of the groups. When applied to a DataFrame, the result is returned as a pandas Series for each column. Examples:

`sum()`
Sum values of each object.

`min()`
Minimum value in each object.

`count()`
Count non-NA/null values of each object.

`max()`
Maximum value in each object.

`median()`
Median value of each object.

`mean()`
Mean value of each object.

`quantile([0.25,0.75])`
Quantiles of each object.

`var()`
Variance of each object.

`apply(function)`
Apply function to each object.

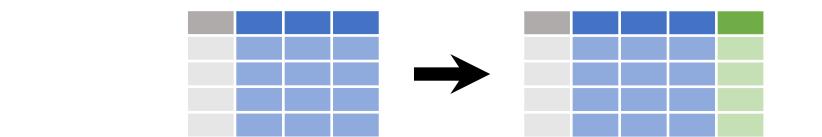
`std()`
Standard deviation of each object.

Handling Missing Data

`df.dropna()`
Drop rows with any column having NA/null data.

`df.fillna(value)`
Replace all NA/null data with value.

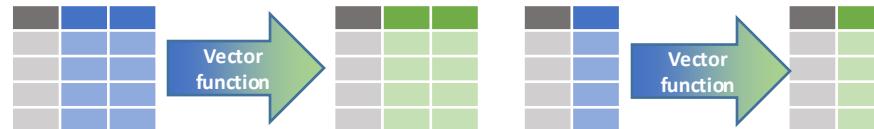
Make New Columns



`df.assign(Area=lambda df: df.Length*df.Height)`
Compute and append one or more new columns.

`df['Volume'] = df.Length*df.Height*df.Depth`
Add single column.

`pd.qcut(df.col, n, labels=False)`
Bin column into n buckets.



pandas provides a large set of **vector functions** that operate on all columns of a DataFrame or a single selected column (a pandas Series). These functions produce vectors of values for each of the columns, or a single Series for the individual Series. Examples:

`max(axis=1)`
Element-wise max.

`min(axis=1)`
Element-wise min.

`clip(lower=-10,upper=10)` `abs()`
Trim values at input thresholds Absolute value.

Windows

`df.expanding()`
Return an Expanding object allowing summary functions to be applied cumulatively.

`df.rolling(n)`
Return a Rolling object allowing summary functions to be applied to windows of length n.

Combine Data Sets

`adf`

x1	x2
A	1
B	2
C	3

`bdf`

x1	x3
A	T
B	F
D	T



Standard Joins

x1	x2	x3
A	1	T
B	2	F
C	3	NaN

`pd.merge(adf, bdf, how='left', on='x1')`
Join matching rows from bdf to adf.

x1	x2	x3
A	1.0	T
B	2.0	F
D	NaN	T

`pd.merge(adf, bdf, how='right', on='x1')`
Join matching rows from adf to bdf.

x1	x2	x3
A	1	T
B	2	F

`pd.merge(adf, bdf, how='inner', on='x1')`
Join data. Retain only rows in both sets.

x1	x2	x3
A	1	T
B	2	F
C	3	NaN
D	NaN	T

x1	x2
A	1
B	2

x1	x2
C	3

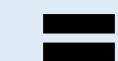
`adf[adf.x1.isin(bdf.x1)]`
All rows in adf that have a match in bdf.

`adf[~adf.x1.isin(bdf.x1)]`
All rows in adf that do not have a match in bdf.

`ydf`

x1	x2
A	1
B	2
C	3

x1	x2
B	2
C	3
D	4



Set-like Operations

x1	x2
B	2
C	3

`pd.merge(ydf, zdf)`
Rows that appear in both ydf and zdf (Intersection).

x1	x2
A	1
B	2
C	3
D	4

`pd.merge(ydf, zdf, how='outer')`
Rows that appear in either or both ydf and zdf (Union).

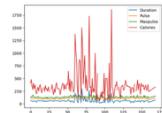
x1	x2
A	1

`pd.merge(ydf, zdf, how='outer', indicator=True)`
`.query('_merge == "left_only"')`
`.drop(columns=['_merge'])`
Rows that appear in ydf but not zdf (Setdiff).

Plotting

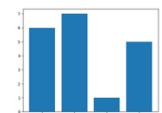
`df.plot()`

Plot a line graph for the DataFrame.



`df.plot.bar()`

Plot a bar graph for the DataFrame.



`df.plot(subplots=True)`

Separate into different graphs for each column in the DataFrame.

`df.plot(title="Graph of A against B")`

Sets the title of the graph.

`df.plot(subplots=True, title=['col1', 'col2', 'col3'])`

Arguments can be combined for more flexibility when graphing, this would plot a separate line graph for of column of a 3-columned DataFrame. The first string in the list of titles applies to the graph of the left-most column.

Changing Type

`pd.to_numeric(data)`

Convert non-numeric types to numeric.

`pd.to_datetime(data)`

Convert non-datetime types to datetime type

`pd.to_timedelta(data)`

Convert non-timedelta types to timedelta

Datetime

With a Series containing data of type datetime, the dt accessor is used to get various components of the datetime values:

`s.dt.year`

Extract the year

`s.dt.month`

Extract the month as an integer.

Mapping

Apply a mapping to every element in a DataFrame or Series, useful for recategorizing or transforming data.

`s.map(lambda x: 2*x)`

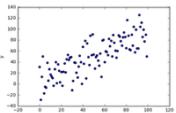
Returns a copy of the series where every entry is doubled

`df.apply(lambda s: s.max() - s.min(), axis=1)`

Returns a Series with the difference of the maximum and minimum values of each row of the DataFrame

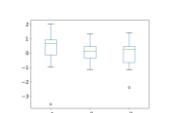
`df.plot.scatter(x='w', y='h')`

Plot a scatter graph of the DataFrame.



`df.plot.boxplot()`

Plot a boxplot of the DataFrame.



`df.plot(cumulative=True)`

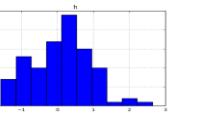
Creates a cumulative plot

`df.plot(bins=30)`

Set the number of bins into which data is grouped (histograms)

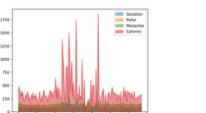
`df.plot.hist()`

Plot a histogram of the DataFrame.



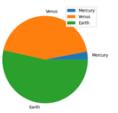
`df.plot.area()`

Plot an area graph of the DataFrame.



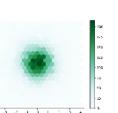
`df.plot.pie()`

Plot a pie chart of the DataFrame.



`df.plot.hexbin()`

Plot a hexbin graph of the DataFrame.



`df.plot(stacked=True)`

Stacks the data for the columns on top of each other. (bar, barh and area only)

`df.plot(alpha=0.5)`

Sets the transparency of the plot to 50%.

Frequently Used Options

Pandas offers some 'options' to globally control how Pandas behaves, display etc. Options can be queried and set via:

`pd.options.option_name` (where `option_name` is the name of an option). For example:

`pd.options.display.max_rows = 20`
Set the `display.max_rows` option to 20.

Functions

`get_option(option)`
Fetch the value of the given option.

`set_option(option)`
Set the value of the given option.

`reset_option(options)`
Reset the values of all given options to default settings.

`describe_option(options)`
Print descriptions of given options.

`option_context(options)`
Execute code with temporary option settings that revert to prior settings after execution.

Display Options

`display.max_rows`
The maximum number of rows displayed in pretty-print.

`display.max_columns`

The maximum number of columns displayed in pretty-print.

`display.expand_frame_repr`

Controls whether the DataFrame representation stretches across pages.

`display.large_repr`

Controls whether a DataFrame that exceeds maximum rows/columns is truncated or summarized

`display.precision`

The output display precision in decimal places.

`display.max_colwidth`

The maximum width of columns, longer cells will be truncated.

`display.max_info_columns`

The maximum number of columns displayed after calling `info()`.

`display.chop_threshold`

Sets the rounding threshold to zero when displaying a Series/DataFrame.

`display.colheader_justify`

Controls how column headers are justified.

Series String Operations

Similar to python string operations, except these are vectorized to apply to the entire Series efficiently.

`s.str.count(pattern)`

Returns a series with the integer counts in each element.

`s.str.get(index)`

Returns a series with the data at the given index for each element.

`s.str.join(sep)`

Returns a series where each element has been concatenated.

`s.str.title()`

Converts the first character of each word to be a capital.

`s.str.len()`

Returns a series with the lengths of each element.

`s.str.cat()`

Concatenate elements into a single string

`s.str.partition(sep)`

Splits the string on the first instance of the separator

`s.str.slice(start, stop, step)`

Slices each string

`s.str.replace(pat, rep)`

Use regex to replace patterns in each string.

`s.str.isalnum()`

Checks whether each element is alpha-numeric

Input/Output

Common file types for data input include CSV, JSON, HTML which are human-readable, while the common output types are usually more optimized for performance and scalability such as feather, parquet and HDF.

`df = pd.read_csv(filepath)`

Read data from csv file

`df = pd.read_html(filepath)`

Read data from html file

`df = pd.read_excel(filepath)`

Read data from xls (and related) files

`df = pd.read_sql(filepath)`

Read data from sql file

`pd.read_clipboard()`

Read text from clipboard

`df.to_parquet(filepath)`

Write data to parquet file

`df.to_feather(filepath)`

Write data to feather file

`df.to_hdf(filepath)`

Write data to HDF file

`df.to_clipboard()`

Copy object to the system clipboard